

Covered: Who Speaks on Cable News, and What Do They Bring?

Gaurav Sood*

August 14, 2026

Journalism’s central convention is attribution. Rather than assert, reporters cite—and the citation is meant to do epistemic work, transferring warrant from a source to a claim (Tuchman 1972). Who gets cited is therefore not merely a question about access. It is a question about what kind of evidence the news rests on.

The sociology of news has long counted sources by their institutional home, and found officials dominant (Sigal 1973; Gans 1979; Bennett 1990). But institutional home and evidentiary standing are different things, and counting only the first conflates them. A prosecutor trying the case being covered and a former prosecutor booked to opine about it occupy the same sector and stand in completely different relations to the facts. The first was there; the second is guessing, expertly. Existing schemes—including recent computational ones that assign each source a single label (Mor-Lan et al. 2025; Spangher et al. 2024)—cannot express that difference.

We measure sources on two crossed axes. *Sector* is the institutional home: government, judiciary, business, academy, media, private individual, and so on. *Epistemic role* is why the source is cited: as the principal whose conduct is the news, as a spokesperson for an institution, as a participant in the events, as an eyewitness, from personal experience, as a credentialed expert, or as a commentator with no privileged access. Normalized mutual

*Gaurav can be reached at: gsood07@gmail.com.

information between the axes is 0.65, so about a third of the epistemic axis is information the sector does not carry.

That third is not spread evenly, and Table 1 is candid about where it lives. In several sectors the epistemic role is close to determined: academics are cited as experts, other outlets’ journalists as commentators, performers as principals, and the second axis buys nothing. The crossing pays where role descriptions vary within a sector—courts and law, private individuals, party operatives, the executive branch—and those are exactly the places where the distinction matters most. It is also where the paper’s claims are concentrated.

The corpus is 6,764,038 external-guest speaker turns from CNN transcripts, 2000–2025, with 530,124 distinct role descriptions. US cable transcripts label speakers as NAME, ROLE:, so the affiliation is supplied in the text rather than inferred, which is what makes classical measurement viable at this scale. Supporting Information SI 1 describes the coding.

What a network plays is not what it books

The clearest result concerns a distinction available only in broadcast. A transcript separates people the network *books* to appear live from people whose recorded words it *plays*, and these are different editorial acts. They are also epistemically different. Figure 1 shows the share of turns from sources with direct access to the events—principals, participants, eyewitnesses, and people speaking from personal experience.

Played clips run 55–74% direct access; live bookings 27–44%. Clips exceed bookings in 25 of 25 years, by a median of 27 percentage points and never by fewer than 18. Clips carry evidence; bookings carry talk. Because this is a within-year comparison, it is unaffected by the measurement drift discussed below.

The asymmetry is easy to rationalize after the fact—a clip is played *because* someone said something consequential, while a booking fills an hour—but it has an uncomfortable implication. The share of airtime given to live panels rather than played material is, in effect,




Figure 1: Share of guest turns from sources with direct access to the events reported, by whether the turn is a live booking or a played clip. Dashed lines mark transcript-format boundaries. The gap holds in every year of the corpus.

a choice about how much of the broadcast rests on evidence.

The executive crowds out everyone else in government

Table 1 gives the joint distribution. Executive-branch sources are 14.8% of guest turns; legislative sources are 1.4%, a ratio of about 11 to one. Cable news covering national politics is overwhelmingly covering the executive.

This is a supply-side counterpart to a demand-side finding. Asked which political leaders come to mind, survey respondents name prominent ones—disproportionately presidents and presidential candidates—rather than ideologically extreme ones, and that paper closed by conjecturing that the bias lies less in who is covered than in how (Sood 2016). On the first half of the conjecture, the supply side is at least consistent with the demand side: the executive branch is what television shows.

Table 1: Institutional sector by epistemic role, CNN guest turns, 2000–2025. Row percentages; % turns is each sector’s share of all guest turns. *Unres.* is reported rather than dropped: it is the share the role text does not settle. H is the normalised entropy of each row: sectors near zero are ones the sector alone already determines, where the second axis adds nothing.

Sector	Prin.	Spok.	Part.	Eyew.	Exp’ce	Expert	Comm.	Pop.	Unres.	% turns	H
Government, executive	69.5	10.3	0.2	0.0	0.0	0.4	19.7	0.0	0.0	14.8	0.38
Government, legislative	78.5	0.2	0.1	0.0	1.5	0.0	19.7	0.0	0.0	1.4	0.27
Government, subnational	65.7	0.5	0.2	0.0	0.6	0.0	33.0	0.1	0.0	0.9	0.33
Judicial and legal	7.0	0.1	16.2	0.0	0.0	5.3	71.3	0.0	0.0	6.1	0.40
Law enforcement	2.8	69.7	5.1	0.0	0.0	22.2	0.1	0.0	0.0	2.5	0.39
Military	5.0	7.8	56.1	0.0	0.4	30.3	0.3	0.1	0.0	1.6	0.49
Party and campaign	37.7	4.7	5.7	0.0	0.0	0.0	51.9	0.0	0.0	3.3	0.46
Business	94.8	0.5	0.0	0.0	0.0	0.1	4.6	0.0	0.0	4.4	0.10
Academic	0.1	0.1	0.0	0.0	0.0	99.7	0.1	0.0	0.0	3.3	0.01
Think tank	5.6	0.0	0.0	0.7	0.0	93.3	0.4	0.0	0.0	0.5	0.13
Nonprofit and advocacy	2.7	93.4	0.0	0.2	0.7	0.7	2.3	0.0	0.0	0.9	0.15
Labor union	25.7	64.8	0.2	0.0	0.1	0.2	9.0	0.0	0.0	0.0	0.40
Religious	0.6	91.3	0.5	0.0	1.4	0.7	5.5	0.1	0.0	0.3	0.18
Media (other outlets)	0.1	0.1	0.0	0.2	0.0	0.1	99.5	0.0	0.0	13.5	0.02
Professional	0.2	0.0	0.1	0.0	0.3	98.7	0.7	0.0	0.0	3.9	0.04
Entertainment and sport	96.4	0.1	0.0	0.0	0.1	0.1	3.4	0.0	0.0	7.4	0.08
Private individual	0.0	0.1	7.7	12.7	71.8	0.0	0.4	7.4	0.0	9.5	0.42
Unknown	0.6	2.5	0.1	0.0	0.1	0.0	0.8	0.1	95.8	25.5	0.10

Ordinary people, and why they are hard to count

Private individuals are 9.5% of guest turns: relatives, neighbours, witnesses, survivors, defendants, refugees, volunteers. This is Gans’s category of “unknowns,” people who reach the news through deviance or misfortune rather than position (Gans 1979).

That figure is a floor, and the reason is itself a finding. How reliably a role description can be classified depends on how often it recurs, because recurring descriptions are standardized job titles and one-off descriptions are not. Among role strings we *can* classify, composition shifts sharply with rarity (Table SI 2.1 in SI 2): private individuals and people speaking from personal experience are far more common among rare strings, and commentators far less. Uncodable strings are rarer than codable ones. So the residue we cannot classify is disproportionately ordinary people speaking from direct experience, and 9.5% understates them.

The mechanism deserves stating plainly. Television labels its pundits in reusable

ways and the people things happen to in bespoke ways—FORMER FBI DIRECTOR recurs; HALEIGH’S BABY-SITTER does not. Any instrument keyed on role descriptions will therefore find the powerful easier to count than the powerless, and will understate the latter unless it says so. Ours currently leaves 25.5% of turns with no sector, and reports that share rather than distributing it.

What we cannot say

We set out to measure whether the evidentiary basis of cable news has changed over twenty-five years, and we cannot answer it with these data. The direct access share is flat across the period, and what moves instead is the split between experts and commentators, which tracks the news cycle rather than time: in 2020 commentary fell and expert testimony rose; in 2024 the reverse. But the more basic problem is that classifiability itself trends upward over the period ($r = 0.82$ with year), because CNN moved from idiosyncratic interview television toward standardized political panels. Worst-case bounds on the direct share are then roughly two and a half times wider than the between-year variation they would have to resolve. Program fixed effects absorb rather more than half the variance in classifiability but under half its trend.

We report this because the alternative—presenting a series whose movement is indistinguishable from the instrument’s improvement—would be worse than reporting nothing.

References

- Asr, F. T., Mazraeh, M., Lopes, A., Gautam, V., Gonzales, J., Rao, P., and Taboada, M. (2021). The gender gap tracker: Using natural language processing to measure gender bias in media. *PLoS ONE*, 16(1):e0245533.
- Bennett, W. L. (1990). Toward a theory of press-state relations in the united states. *Journal of Communication*, 40(2):103–125.
- Gans, H. J. (1979). *Deciding What’s News: A Study of CBS Evening News, NBC Nightly News, Newsweek and Time*. Pantheon, New York.
- Mor-Lan, G., Sheaffer, T., and Shenhav, S. R. (2025). You’re not gonna believe this: A computational analysis of factual appeals and sourcing in partisan news. *arXiv preprint arXiv:2510.10658*.
- Sigal, L. V. (1973). *Reporters and Officials: The Organization and Politics of Newsmaking*. D. C. Heath, Lexington, MA.
- Sood, G. (2016). Extreme recall: Which politicians come to mind? Working paper.
- Spangher, A., Youn, J., DeButts, M., Peng, N., Ferrara, E., and May, J. (2024). Explaining mixtures of sources in news articles. In *Findings of the Association for Computational Linguistics: EMNLP 2024*.
- Tuchman, G. (1972). Objectivity as strategic ritual: An examination of newsmen’s notions of objectivity. *American Journal of Sociology*, 77(4):660–679.

Supporting Information

SI 1 The taxonomy

Each speaker turn is parsed from the `NAME, ROLE:` label. A role clause is written only at a speaker’s first mention; later turns are bare surnames, so roles are propagated within a transcript from a roster, and whether a role was read locally or inherited is recorded.

Two independent lexicons map keywords to sector and to epistemic role. Matching is longest-keyword-wins on whole words, so row order in the lexicon files is irrelevant and a new keyword cannot shadow an existing one. Because the axes are matched separately, they compose: `PENTAGON SPOKESMAN` resolves to `military × spokesperson` without either file enumerating the pair.

Three attributes are lifted out of the role clause before matching, because they describe delivery or affiliation rather than the kind of source: how the speaker reached air (*voice-over, via telephone, through translator*), a party tag, and standing (*current, former, candidate*). Standing then modifies the epistemic role: leaving a post ends the standing that made someone a principal, a spokesperson or a participant, so a former officeholder is coded as a commentator. Where authority is a credential rather than a post, whether the source is an expert or a commentator depends on the segment topic, and those cells are flagged rather than guessed.

When no keyword matches, three fallbacks apply in order: a typographic convention by which an outlet name alone denotes a visiting journalist; an honorific in the speaker’s name (US legislators are written `SEN. JOHN MCCAIN, (R) ARIZONA`, putting the office in the name and only party and state in the role); and a frozen per-string dictionary for descriptions no rule generalizes to. Every label records which route produced it.

Table SI 2.1: Epistemic composition of codable role strings, by how often the string recurs. Quintiles of string frequency; percentages within quintile. Categories over-represented among rare strings are the ones the uncoded residual is mostly made of, and are therefore under-counted in Table 1.

Epistemic role	Rarest %	Commonest %	Gap (pp)
Comm.	16.5	33.0	-16.5
Prin.	34.1	34.6	-0.5
Pop.	0.8	1.0	-0.2
Expert	11.5	11.5	-0.1
Eyew.	3.1	1.5	+1.6
Part.	6.1	3.8	+2.2
Spok.	10.2	6.1	+4.2
Exp'ce	17.6	8.5	+9.1

SI 2 Coverage and what is missing

Table SI 2.1 reports epistemic composition of classifiable role strings by quintile of how often the string recurs. Read the gap column as the sign of the bias in Table 1: categories over-represented among rare strings are what the unclassified residue is mostly made of, and are therefore undercounted.

Epistemic roles are read from the text for 71.6% of turns; the remainder fall back on what the sector implies, and are marked as such throughout so that any table can be recomputed without them.

One consequence deserves stating. Some keywords fix both axes at once—a professor is academic and an expert—so for sectors whose role vocabulary is homogeneous, the epistemic label is partly an artifact of how the lexicon is built rather than an independent reading of the text. The *H* column in Table 1 makes those rows visible, and the aggregate mutual information should be read as an average over sectors that differ a great deal in how much the second axis is doing.

SI 3 Validation

[**VALIDATION PENDING**] A stratified probability sample of speaker turns, double-coded, will supply: inter-coder agreement per axis; a confusion matrix per axis rather than a headline accuracy; and recall by sector and by epistemic role separately, since differential recall—not average accuracy—is what biases composition estimates ([Asr et al. 2021](#)). An external check is available at no cost: a sister corpus of MSNBC transcripts ships a human-authored guest list for every transcript, sharing no code with this pipeline.